

Re-Identification in a single feed

1. Overview & Pipeline

We built a robust, end-to-end re-identification system that assigns consistent IDs to players—even under occlusion, collisions, and entry/exit—in a 15-second clip.

1. Detection

- Fine-tuned YOLOv11 (via Ultralytics code) to output only the “player” class.
- Confidence threshold 0.25, IoU filter 0.5.

2. Feature Extraction

- Deep CNN (MobileNetV2) backbone from PyTorch for 1,280-D appearance embeddings.
- Crops are resized to 128×256, then L2-normalized.

3. Jersey OCR

- Preprocess each crop (CLAHE, sharpening, adaptive threshold).
- Tesseract in single-char mode (`--psm 10`) with digit-whitelist.
- Track-level history of 5 frames; jersey confirmed once all 5 match.

4. Tracking & Data Association

- **Kalman filter** (8-state constant-velocity) predicts each track’s bounding box and speed.
- **Cost matrix** combines three terms:
 - $0.4 \times (1 - \text{IoU}(\text{predicted}, \text{detected}))$
 - $0.4 \times (1 - \text{cosine}(\text{appearance}))$

- $0.2 \times$ (jersey-match penalty: 0 if same, 1 otherwise)
- **Motion gating**: any detection whose center jumps beyond $1.5 \times$ predicted speed (or a 30-px minimum) is disallowed.
- **Hungarian assignment** on gated cost matrix; unmatched detections spawn *tentative* tracks.

5. Track Lifecycle Management

- **Confirmation**: tracks require 3 consecutive hits before moving from TENTATIVE to CONFIRMED.
- **Survival**: confirmed tracks remain alive for up to 50 missed frames before deletion.
- **Output filtering**: only report tracks that have been confirmed *and* persisted for at least 5 frames, eliminating ghost IDs.
- **Dead-track cooldown**: after a track is deleted, its ID is barred from being reassigned for 50 frames, preventing immediate ID reuse.

2. Experimental Techniques & Rationale

- **HSV histograms → CNN embeddings**
Switched from 96-D color histograms to 1,280-D MobileNetV2 features. Learned embeddings vastly improve discrimination under similar kits and partial occlusions.
- **Single-char OCR + temporal smoothing**
Running Tesseract in `--psm 10` mode on well-preprocessed crops gave cleaner digit reads. Holding a 5-frame majority vote stabilizes jersey assignments.
- **Motion gating**
Filters out implausible associations when two players collide, leveraging the Kalman-estimated velocity. This dramatically cuts swap rates during overlaps.
- **Output gating & cooldown**
Delaying track outputs until stable prevents showing fleeting, noisy IDs. The dead-track cooldown stops brand-new players from inheriting old IDs.

3. Remaining Challenges & Future Work

- **Full occlusions** still occasionally cause ID gaps when a player lingers off-frame longer than the cooldown.
- **Heavy overlaps** in dense scenes can overwhelm motion gating.
- **OCR failure** on extreme motion blur.

Next steps could include:

- Integrating a specialized person re-ID network (e.g. OSNet) trained on player datasets.
- Adding a “merge-and-split” handler for two-into-one collision blobs.
- Extending the cooldown logic to incorporate appearance confidence for smarter ID revival.